

Measured trajectories (real 4-model deployment)


Reputation Update Rule (Algorithm 1)

Correct vote (ratio < 0.5): +0.1

Honest mistake (ratio < 0.7): -0.05

Persistent rejection (ratio ≥ 0.7): -0.25

Correct vote (ratio 0.5-0.7): -0.02



Reputation updates (chronological, across tasks)

—●— honest — GSM8K $n=8, f=2, s=1$

—■— honest — MBPP $n=8, f=2, s=1$

—●— Byzantine — GSM8K $n=8, f=2, s=1$

—■— Byzantine — MBPP $n=8, f=2, s=1$